

Attributing annotator polarization to sociodemographic groups

Dimitris Tsirmpas, John Pavlopoulos

March 2025

1 Introduction

Annotations are essential for a wide range of tasks, especially in fields such as content moderation, toxicity detection and hate speech detection. These tasks are especially essential for training systems that protect vulnerable and minority groups in online spaces, where they are often targeted [17, 4]. However, annotations for such tasks are subjective, and are usually based on majority-group annotators. This may limit the usefulness of the data, or even lead to biased and misconfigured systems.

The core of the problem is simple; if a comment targets marginalized groups, there can be cases where most annotators overlook it, while the few marginalized annotators vehemently disagree in vain. An example of such cases would be racist comments that are presented with coded language (usually referred to as “dog-whistles” [13]), which are not picked up by most annotators, but only by ones belonging in the targeted minority. Inversely, majority-group annotators may bias their own annotations against the minority-group annotators; CITATION for instance shows that toxicity models disproportionately marked African American speech as “toxic”. This issue exists even in representative samples.

Inter-annotator agreement is often used to detect such cases [8]. Metrics based on agreement mostly measure annotation variance, which may not be an optimal problem formulation for detecting such cases. *Polarization* is a better instrument in this regard, since it formulates annotation analysis as a cluster identification problem [11, 12]—specifically whether two or more annotation clusters are present. Figure 1 displays how the two methods handle four common cases of annotation distributions for a hypothetical




Figure 1: A comment annotated by 100 annotators. Disagreement is similar to variance; thus it overlooks cases where annotators may be “split down the middle”. Polarization on the other hand, is able to identify multiple clusters, even if they are close together.

toxicity annotation task.

While we can detect polarization [12, 11], there is little work on how to detect whether it is actually caused by marginalized groups. Pavlopoulos and Likas [12] propose a mechanism which only works in the extreme case where overall polarization is zero; in this work we demonstrate that this case is not frequent in real-world data. We instead propose the “*Aposteriori Unimodality (Apunim) Metric*”, which attributes polarization to individual annotator sociodemographic groups. Our contributions are:

1. We introduce a new metric which measures how much polarization can be attributed specific sociodemographic groups.




Figure 2: Example of a polarizing comment, where male and female annotators agree between themselves, but disagree with the opposite gender.

2. We discover issues not yet acknowledged in literature, such as the presence of inherent polarization and conflicting polarization directions. Our proposed metric is designed to avoid these issues.
3. We apply our metric to four datasets; two with human comments and annotations [15, 9] as well as two generated and annotated by Large Language Model (LLM) agents [16].

2 Methodology

In this section, we provide a formal mathematical formulation for the problem of attributing polarization to specific annotator characteristics (§2.1), and offer an intuitive rationale for how established polarization metrics can be leveraged (§2.2). We then explain how we can measure polarization for each group (§2.3) and how much of that polarization is caused by external variables (§2.4). Finally, we introduce our metric, which computes attribution on the entire dataset for each sociodemographic group (§2.5).

2.1 Problem Formulation

SocioDemographic Backgrounds Since our goal is to pinpoint which specific characteristics contribute to polarization, we need to isolate individual groups within a SocioDemographic Background (SDB). Let Θ be one of multiple “dimensions” the set of all annotator SDB groups for a single dimension (e.g. “male”, “female” and “non-binary” if we are investigating annotator gender).

Annotations Let $d = \{c_1, c_2, \dots\}$ be a dataset composed of multiple annotated data points. We assume that annotating a data point depends on two variables: inherent (“apriori”) controversy, and the (“aposteriori”) annotator’s SDB, as well as uncontrolled factors such as mood and personal experiences (noise). Assuming that each data-point c is assigned multiple annotators, we can define the annotation multi-set $A(c) = \{(a, \theta)\}$, where a is a single annotation. As an example, the annotations for a comment in a sentiment analysis task would be formulated as:

$$A(\text{"Could be better, could be worse"}) = \{(\text{positive}, \text{female}), (\text{negative}, \text{male}), (\text{positive}, \text{female}), \dots\} \quad (1)$$

Polarization Each set of annotations features a certain degree of *polarization*. Unpolarized annotations are usually unimodal; the annotators disagree on the details (which would be shown roughly as a bell curve around the median annotation), not fundamentally (which would likely be shown as a multimodal distribution). Pavlopoulos and Likas [12] create a polarization metric, the “normalized Distance From Unimodality (nDFU)”, which measures whether an annotation set shows no polarization (unimodal distribution— $nDFU \approx 0$), up to complete polarization (multimodal distribution— $nDFU \approx 1$). Thus, we need a metric that measures how much the nDFU of a polarizing data point can be partially explained by θ .




Figure 3: Hypothetical example of a polarizing discussion with two comments, for which the annotators disagree on which is the toxic one. If we aggregate the two comments, the polarization scores for both men and women significantly rise, obscuring whether the exhibited polarization can be partially attributed to gender.

2.2 Quantifying changes in polarization

Comment polarization Intuitively, θ partially explains the polarization of a data point c when the annotations grouped by θ are more polarized compared to the full set of annotations. Figure 2 exhibits a hypothetical example where a misogynistic comment is annotated for toxicity by male and female annotators.¹ The annotations are generally polarized ($nDFU_{all} = 0.625$). The annotations by female annotators exhibit low polarization between themselves ($nDFU_{women} = 0.1$), since most agree the data-point is toxic. The set of male annotations, also shows low polarization ($nDFU_{men} = 0.3725$), but for the opposite reason—most men agree that it is *not* toxic. This suggests that the overall polarization is driven by disagreements between male and female annotators.

Dataset polarization Given this observation, we would be tempted to aggregate all annotations in the dataset and compare the polarization of each θ with the full set of annotations $A(c)$. In fact this is very tempting since aggregating annotations solves the very real and important problem of annotation scarcity; each data point may be annotated by only

5 annotators, leading to a best-case 3-2 split between groups. Since $nDFU$ is at its core a histogram estimation, 3 annotations may not be enough to acquire a reliable histogram (this is discussed in depth in §2.5).

However, this formulation would not work well. Figure 3 illustrates a hypothetical discussion with two comments, both of which are toxic, but where male and female annotators disagree on *which* comment is the toxic one. If we aggregate the annotations to a single discussion, the opposing polarization effects might balance each other out, leading to a false negative. This example demonstrates that polarization has a *direction*, which is not captured by the (symmetric) $nDFU$ metric.

2.3 Observed polarization

Definition The observed polarization of a data-point on the group of annotations θ can be directly computed as:

$$pol(c, \theta) = nDFU(P(A(c), \theta)) \quad (2)$$

where $P(A(c), \theta_i) = \{a(c; \theta) \in A(c) | \theta = \theta_i\}$ is the partition of A for a data-point c , for which its annotators belong to the SDB group θ_i .

How many annotators? While our test works best with many annotations per datapoint, the cost

¹The use of only two predominant genders is made only for the purposes of demonstration.




Figure 4: Standard error of the *pol-statistic* when sampled with replacement from the 100 LLM-annotator synthetic dataset for various number of annotators.

required for multiple human annotators is often prohibitive [14]. Figure 4 demonstrates that our metric becomes consistent with about 15 annotators per comment. This experiment was conducted on a subset of the Sap et al. [15] dataset, with LLM annotators having different SDBs. Since we use LLM annotators on one dataset, we can not claim that this finding generalizes to human annotation; it serves only as a preliminary indication. The full experimental procedure can be found in Appendix D.

2.4 Apriori polarization

Definition In §1 we mentioned that annotator polarization can be caused either by the annotators SDB, or other factors. The latter can be considered “apriori” polarization, which is driven by the inherent “controversy” of the data point. Previous work [12] considered the case where this apriori polarization was zero; this is not the case in many datasets, as we demonstrate in §3.

Statistical problems In §2.2 we compared the polarization of the whole set of annotations with the partitions by annotator group. This approach uses the same samples for both subset and superset, violating the fundamental hypothesis of independent




Figure 5: Since we can not compare the statistics of the group-subsets with the annotation superset, we elect to instead compare them with the mean apriori polarization of many sampled subsets.

samples. This also means the statistics for both will *always be similar to an extent*. This is especially pronounced in our case, due to the extreme overlap exhibited by tiny annotation sets. Additionally, when calculating p-values, the non-independence of samples induces an intrinsic, non-zero covariance between any statistic s calculated on A and on S .

A standard method of side-stepping this issue is to compare the subset with its complement to the superset. This can not be applied to our case. Consider the example in Figure 2; if we compared $nDFU_{men}$ with $nDFU_{women}$, the result would be very close to zero, indicating no polarization attribution to gender.

Overcoming dependent samples Given infinite annotations, the assumption of independence would hold. Likewise, if we compared the subset with infinite random subsets of the superset. We can thus estimate the apriori polarization in data point c by bootstrapping (Figure 5). We randomly partition the annotations in $|\Theta|$ groups, with matching group sizes (e.g., if we have 100 annotations, 80 of which are made by male annotators and 20 by female annotators, we will create random partitions of sizes 80 and




Figure 6: An overview of the Aposteriori Unimodality Test. We gather statistical information from the annotations (different SDBs are denoted by different colors) via the $nDFU$ measure, which is aggregated on the data-point-level by the pol -statistic. The Aposteriori Unimodality Test is applied on the discussion level, operating on the pol -statistics of the individual data-points.

20). The randomly partitioned annotations are then sampled t times. Formally, the estimated apriori polarization will be given by:

$$\frac{1}{t} \sum_{i=1}^t nDFU(\tilde{P}_i(A(c))) \quad (3)$$

where $\tilde{P}_i$ is the random partition operator with matching group sizes.

2.5 Aposteriori polarization

Dataset filtering While data scarcity is an important limitation when acquiring the observed polarization, this is decisively not the case with the number of data points themselves, which can range up to hundreds of thousands. Thus, we have the luxury of filtering data points before computing a dataset-wide metric. Since our metric attributes polarization, unpolarized data points are useless. The same applies for data points which are composed of only a single annotator group. Thus, we define the subset S_d of dataset d such that:

$$S_d = \{c \in d \mid nDFU(c) > \alpha \text{ and } c \text{ annotated by at least two groups}\} \quad (4)$$

where α is a small positive number (in our experiments $\alpha = 0.2$).

Definition By obtaining the pol -statistics for all data-points in S_d for SDB θ , we can get the mean observed polarization for the dataset:

$$pol_{obs.}(d, \theta) = \frac{1}{|d|} \sum_{c \in S_d} pol(c, \theta) \quad (5)$$

and similarly the mean apriori dataset polarization:

$$pol_{apriori}(d) = \frac{1}{|d|} \sum_{c \in S_d} \frac{1}{t} \sum_{i=1}^t nDFU(\tilde{P}_i(A(c))) \quad (6)$$

We can then define the “*Aposteriori Unimodality*” (*Apunim*) Statistic as:

$$apunim(d, \theta) = \frac{pol_{obs.}(d, \theta) - pol_{apriori}(d)}{1 - pol_{apriori}(d)} \quad (7)$$

Since $nDFU \rightarrow [0, 1] \Rightarrow apunim \rightarrow [-1, 1]$. If $apunim(d, \theta) \approx 0$, the exhibited polarization among annotators of group θ does not surpass what would be expected by chance. If $apunim(d, \theta) > 0$, then the group partially explains a rise in polarization, and inversely for the case $apunim(d, \theta) < 0$. The scope of the statistic and its relationship with the previously presented statistics are demonstrated in Figure 6.

p-value estimation While it is possible to obtain p-value results for the $apunim$ values being larger or less than 0 respectively, we find that the inherent variance of our statistic makes a p-value estimation worthless ($p \gg 0.1$). Nevertheless, we include a definition for parametric and non-parametric p-value estimation, and apply them to our datasets in Appendix B.

3 Datasets

We use the datasets provided by Kumar et al. [9] and Sap et al. [15]. The datasets feature various online comments, each annotated by a number of human annotators (5 and 4-6 annotators per comment




Figure 7: Histogram of dataset polarization for the human datasets considered in this paper per SDB feature.

respectively). There is a reasonable amount of polarization in most annotations for both dataset, as demonstrated by Figure 7.

The Sap et al. [15] dataset includes racism annotations for 626 Twitter/X comments, and standard SDB information (age, race, education, gender). The Kumar et al. [9] dataset measures toxicity on various online comments, and the provided SDBs include mostly annotator experiences (e.g., whether they have been personally targeted online) as well as standard SDB information such as sexual orientation, age and education. Since the dataset is extensive (106035 comments), we only select a sample of 2000, due to computational constraints. Further analysis of the datasets can be found in Appendix TODO.

4 Results

The results for the Kumar et al. [9] and Sap et al. [15] datasets can be found in Tables 1 and 2 respectively.

The only shared demographic dimension between the two datasets is *Age*, which consistently shows low apunim scores. The [15] dataset shows large polarization attribution to *Ethnicity* and *Gender*, relative to *Age*. This is an expected result given our intuition §1.

Results for the [9] dataset are less conclusive, with no apunim values exceeding 0.1 in any direction. In relative terms, the annotator being *Parent*, *Transgen-*

Table 1: Apunim results for the Sap et al. 2022 dataset

SDB Feature		apunim
Age	0-20	—
	20-40	0.0017
	40-60	0.0115
	60-80	0.0049
Ethnicity	black	-0.2967
	hisp	—
	middleEastern	—
	other	—
	white	0.1400
Gender	man	0.0942
	nonBinary	—
	woman	-0.1810

der and *Religious* seems to be the most important dimensions when it comes to polarization attribution. There is also a notable deviation between *Heterosexual* and *LGBT* annotators, which is not present between different groups of *LGBT* annotators. We also observe a clear rising polarization attribution the more *Religious* the annotators become.

Table 2: Apunim results for the Kumar et al. 2021 dataset

SDB Feature		apunim
Age	18 - 24	-0.0120
	25 - 34	0.0267
	35 - 44	-0.0252
	45 - 54	0.0065
	55 - 64	-0.0098
	65 or older	-0.0040
	Prefer not to say	0.0000
	Under 18	0.0000
Education	Associate degree	-0.0046
	Bachelor's degree	0.0229
	College, no degree	-0.0424
	Doctoral degree	0.0053
	High School graduate	-0.0148
	Master's degree	0.0162
	No high school	-0.0147
	Other	0.0000
	Prefer not to say	0.0000
	Professional degree	-0.0056
Has Been Targeted	False	-0.0422
	True	0.0047
Is Parent	No	-0.0813
	Prefer not to say	0.0133
	Yes	0.0403
Is Transgender	No	-0.0732
	Prefer not to say	0.0084
	Yes	-0.0113
Political Affiliation	Conservative	0.0476
	Independent	-0.0142
	Liberal	-0.0328
	Other	-0.0218
	Prefer not to say	0.0025
Seen Toxicity	False	0.0451
	True	-0.0898
Sexual Orientation	Bisexual	0.0067
	Heterosexual	-0.0522
	Homosexual	-0.0098
	Other	0.0016
	Prefer not to say	-0.0062
	NaN	0.0210
Religion Is Important	Not important	-0.0647
	Not too important	-0.0153
	Prefer not to say	-0.0018
	Somewhat important	0.0094
	Very important	0.0442
Toxicity Is Problem	Frequently a problem	-0.0156
	Not a problem	0.0009
	Occasionally a problem	-0.0180
	Rarely a problem	0.0097
	Very frequently a problem	-0.0199

5 Conclusion

We introduced the ‘‘Aposteriori Unimodality Test’’, a statistical test that detects whether certain sociodemographic groups partly cause polarization in annotations. Our test is resistant to low samples of annotations, bypasses common statistical issues such as

sample dependence, and avoids newly-discovered issues such as the presence of inherent polarization and polarization direction. We apply our test to two human-generated and two LLM-generated datasets and find interesting patterns on the former, and verify current literature on LLM SDB prompting on the latter.

6 Limitations

Since there are no other established tests that attribute polarization to sociodemographic characteristics, there is no way to verify that the results presented in §4 are accurate. Similarly, there is no qualitative way of evaluating our test, other than observations on the (non-) efficacy of LLM SDB prompting and our own intuition. Lastly, while our test seems to work sufficiently well with a small number of annotators per data-point, we still encourage researchers using this test to have at least a few annotators of each sociodemographic group that is being tested, in their dataset.

7 Ethical Considerations

While our work aims to help protect marginalized and disadvantaged groups by attributing polarization to certain subgroups, it can be taken advantage of by malicious actors. Given a dataset with fine-grained SDB information, these actors can use our test to target specific vulnerable groups. We thus urge researchers to keep datasets including such information protected, and provided to others only under explicit permission.

References

- [1] Jacy Reese Anthis et al. *LLM Social Simulations Are a Promising Research Method*. 2025. arXiv: 2504.02234 [cs.HC]. URL: <https://arxiv.org/abs/2504.02234>.
- [2] James Bisbee et al. ‘‘Synthetic Replacements for Human Survey Data? The Perils of Large Language Models’’. In: *Political Analysis* 32.4 (2024), 401–416. DOI: 10.1017/pan.2024.5.

- [3] S.Y. Chen, Z. Feng, and X. Yi. “A general introduction to adjustment for multiple comparisons”. In: *Journal of Thoracic Disease* 9.6 (2017). doi: 10.21037/jtd.2017.05.34; PMID: 28740688; PMCID: PMC5506159, pp. 1725–1729.
- [4] Shelby Dioum. *What Explains the Increase in Online Hate Speech?* Accessed: 2025-09-22. University of California, Davis, Mar. 2025. URL: <https://www.ucdavis.edu/magazine/what-explains-increase-online-hate-speech>.
- [5] Luke Hewitt et al. “Predicting Results of Social Science Experiments Using Large Language Models”. Equal contribution, order randomized. Aug. 2024.
- [6] Sture Holm. “A Simple Sequentially Rejective Multiple Test Procedure”. In: *Scandinavian Journal of Statistics* 6.2 (1979), pp. 65–70. ISSN: 03036898, 14679469. URL: <http://www.jstor.org/stable/4615733> (visited on 07/07/2025).
- [7] Bernard J. Jansen, Soon gyo Jung, and Joni Salminen. “Employing large language models in survey research”. In: *Natural Language Processing Journal* 4 (2023), p. 100020. ISSN: 2949-7191. DOI: <https://doi.org/10.1016/j.nlp.2023.100020>. URL: <https://www.sciencedirect.com/science/article/pii/S2949719123000171>.
- [8] Petra Kralj Novak et al. “Handling Disagreement in Hate Speech Modelling”. In: *Information Processing and Management of Uncertainty in Knowledge-Based Systems*. Ed. by Davide Ciucci et al. Cham: Springer International Publishing, 2022, pp. 681–695. ISBN: 978-3-031-08974-9.
- [9] Deepak Kumar et al. “Designing toxic content classification for a diversity of perspectives”. In: *Proceedings of the Seventeenth USENIX Conference on Usable Privacy and Security*. SOUPS’21. USA: USENIX Association, 2021. ISBN: 978-1-939133-25-0.
- [10] Terrence Neumann, Maria De-Arteaga, and Sina Fazelpour. *Should you use LLMs to simulate opinions? Quality checks for early-stage deliberation*. 2025. arXiv: 2504.08954 [cs.CY]. URL: <https://arxiv.org/abs/2504.08954>.
- [11] John Pavlopoulos and Aristidis Likas. “Distance from Unimodality for the Assessment of Opinion Polarization”. In: *Cognitive Computation* 15.2 (2023), pp. 731–738. ISSN: 1866-9964. DOI: 10.1007/s12559-022-10088-2. URL: <https://doi.org/10.1007/s12559-022-10088-2>.
- [12] John Pavlopoulos and Aristidis Likas. “Polarized Opinion Detection Improves the Detection of Toxic Language”. In: *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Yvette Graham and Matthew Purver. St. Julian’s, Malta: Association for Computational Linguistics, Mar. 2024, pp. 1946–1958. URL: <https://aclanthology.org/2024.eacl-long.117/>.
- [13] Anne Quaranto. “Dog whistles, covertly coded speech, and the practices that enable them”. In: *Synthese* 200.4 (2022), p. 330.
- [14] Luca Rossi, Katherine Harrison, and Irina Shklovski. “The Problems of LLM-generated Data in Social Science Research”. In: *Sociologica* 18.2 (2024), 145–168. DOI: 10.6092/issn.1971-8853/19576. URL: <https://sociologica.unibo.it/article/view/19576>.
- [15] Maarten Sap et al. “Annotators with Attitudes: How Annotator Beliefs And Identities Bias Toxic Language Detection”. In: *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Ed. by Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz. Seattle, United States: Association for Computational Linguistics, July 2022, pp. 5884–5906. DOI: 10.18653/v1/2022.naacl-main.431.

URL: <https://aclanthology.org/2022-naacl-main.431/>.

- [16] Dimitris Tsirmpas, Ion Androutsopoulos, and John Pavlopoulos. *Scalable Evaluation of Online Facilitation Strategies via Synthetic Simulation of Discussions*. 2025. arXiv: 2503.16505 [cs.HC]. URL: <https://arxiv.org/abs/2503.16505>.
- [17] United Nations. *Targets of Hate*. Accessed: 2025-09-22. Sept. 2025. URL: <https://www.un.org/en/hate-speech/impact-and-prevention/targets-of-hate>.

A Acronyms

VMD	Virtual Moderation Dataset
LLM	Large Language Model
SDB	SocioDemographic Background
nDFU	normalized Distance From Unimodality
FWER	Family-Wise Error Rate

B p-value Estimation

B.1 Methodology

Like most metrics, we also include a p-value alongside the *apunim* value. This can be computed either non-parametrically, by repeatedly shuffling annotations and computing a null distribution of *apunim*, or parametrically, by assuming a normal distribution (which is generally a safe assumption if $t \geq 30$). In the latter case, we compute:

$$z = \frac{\hat{\kappa} - \mathbb{E}[\kappa]}{SE(\hat{\kappa})}, p = 1 - \Phi(z) \quad (8)$$

Since we are simultaneously considering $|\Theta|$ hypotheses, we apply a multiple comparison correction to the resulting p-values. We choose the Holms method [6]. The test is parameterized by the Family-Wise Error Rate (FWER), which is used to tune the strength of the correction; we can increase this

value to make our test more conservative towards multiple hypotheses [3]). In general, it is safe to set FWER equal to the significance level of our test (e.g., $FWER = 0.95$ if we are looking for $p < 0.05$).

B.2 Results

The full results for the datasets presented in §3, including p-value estimations, can be found in Tables 4 and 3 for the Sap et al. [15] and Kumar et al. [9] datasets respectively. The experiments were ran on the exact same parameters as in §4.

Note the absence of any statistically significant results for both parametric and non-parametric tests. This suggests that, although the differences in polarization captured by the *pol* statistic introduced in §2 are substantial enough to hint at polarization being attributable to certain groups, they remain too small to overcome the inherent noise in human annotations. As a result, we cannot obtain a robust estimate of statistical significance.

C Synthetic Datasets

We use two synthetic datasets; one is the Virtual Moderation Dataset (VMD) presented in Tsirmpas, Androutsopoulos, and Pavlopoulos [16]. This dataset features 140 discussions, each having 14 comments (and usually up to 14 facilitator comments), each comment annotated by 10 LLM annotators, each supplied with a different SDB. The second dataset, is an extension of VMD, where we select four random unmoderated discussions, and employ 100 distinct LLM annotators.

An issue with synthetic datasets is the lack of overall polarization (see Figure 8). This can be offset by the ability of synthetic datasets to be scaled to a much greater extent than human datasets. This is because LLMs do not have cost and availability constraints of humans.

The results for the 100-annotator dataset can be found in Table ???. We can not obtain results for the VMD dataset, since no comment exhibited apriori polarization while also being represented by more than two annotators of different groups. The lack




Figure 8: Histogram of dataset polarization for the synthetic datasets considered in this paper per SDB feature.

of overall polarization, as well as the quantitatively smaller apunim values, suggest that SDB prompting can not be used to simulate annotations for human groups. This is consistent with relevant literature on the topic [1, 5, 14, 7, 2, 10].

D LLM Annotation

TODO

Table 3: Apunim results for the kumar dataset

SDB Feature		apunim	p-param	p_nonparam
Age	18 - 24	-0.0437	0.8916	1.0000
	25 - 34	0.0589	0.9199	0.9750
	35 - 44	-0.0397	0.9507	0.9677
	45 - 54	0.0117	0.9602	1.0000
	55 - 64	0.0179	0.9424	1.0000
	65 or older	—	—	—
	Prefer not to say	—	—	—
Education	Associate degree	-0.0257	0.9306	1.0000
	Bachelor’s degree	-0.0611	0.9237	0.9024
	College, no degree	-0.0763	0.8670	0.9455
	Doctoral degree	—	—	—
	High School graduate	-0.0627	0.8811	1.0000
	Master’s degree	0.0944	0.8265	1.0000
	No high school	-0.2766	0.7920	1.0000
	Other	—	—	—
Has Been Targeted	Professional degree	—	—	—
	False	-0.0153	0.9824	1.0000
Is Parent	True	-0.0261	0.9701	0.9583
	No	-0.1724	0.7727	0.8293
	Prefer not to say	—	—	—
Is Transgender	Yes	0.1527	0.7951	0.9024
	No	-0.1105	0.8654	0.8824
	Prefer not to say	—	—	—
Political Affiliation	Yes	-0.0695	0.8842	1.0000
	Conservative	0.0316	0.9471	0.9859
	Independent	-0.0026	0.9968	1.0000
	Liberal	0.0197	0.9681	0.9398
	Other	-0.0989	0.9565	1.0000
	Prefer not to say	0.0044	0.9744	1.0000
Seen Toxicity	False	-0.0275	0.9679	1.0000
	True	-0.0574	0.9103	0.9701
Sexual Orientation	Bisexual	0.0337	0.9177	1.0000
	Heterosexual	-0.1301	0.7809	0.7818
	Homosexual	0.0280	0.8207	1.0000
	Other	—	—	—
	Prefer not to say	0.0123	0.9417	1.0000
Thinks Religion Is Important	Not important	-0.1930	0.7213	0.8060
	Not too important	-0.0406	0.9261	1.0000
	Prefer not to say	—	—	—
	Somewhat important	0.0033	0.9890	1.0000
	Very important	0.0545	0.9127	0.9726
Thinks Toxicity Is Problem	Frequently a problem	-0.1537	0.7922	0.8831
	Not a problem	-0.0384	0.8425	1.0000
	Occasionally a problem	0.0876	0.8531	0.8974
	Rarely a problem	-0.0152	0.9775	1.0000
	Very frequently a problem	0.0162	0.9537	1.0000

Table 4: Apunim results for the sap dataset

SDB Feature		apunim	p_param	p_nonparam
Age	0-20	—	—	—
	20-40	0.0017	0.9958	0.9977
	40-60	0.0115	0.9794	0.9976
	60-80	0.0049	0.9660	1.0000
Ethnicity	black	-0.2967	0.6369	0.6564
	hisp	—	—	—
	middleEastern	—	—	—
	other	—	—	—
	white	0.1400	0.7130	0.8246
Gender	man	0.0942	0.8368	0.8920
	nonBinary	—	—	—
	woman	-0.1810	0.7198	0.7653

Table 5: Apunim results for the 100 dataset

SDB Feature		apunim	p_param	p_nonparam
age	0-20	0.0447	0.8422	0.8256
	20-40	-0.0310	0.9492	0.9590
	40-60	-0.0829	0.8636	0.8667
	60-80	0.0415	0.9300	0.9385
education_level	bachelor’s degree	-0.0649	0.8474	0.8564
	doctoral degree	-0.0107	0.9614	0.9744
	master’s degree	-0.0155	0.9609	0.9538
	no formal education	-0.0596	0.8557	0.8769
	primary education	0.0275	0.9392	0.9487
	secondary education	0.0197	0.9676	0.9846
	vocational training	0.0007	0.9970	1.0000
gender	female	0.0095	0.9833	1.0000
	male	-0.0144	0.9807	0.9795
	non-binary	0.0298	0.8955	0.9231
political_affiliation	apolitical	-0.0381	0.9238	0.9128
	centrist	-0.0100	0.9851	0.9795
	left-wing	-0.0205	0.9670	0.9641
	right-wing	0.0044	0.9899	0.9897
sexual_orientation	asexual	0.0031	0.9876	0.9949
	bisexual	-0.0338	0.9017	0.9077
	heterosexual	-0.0037	0.9973	0.9846
	homosexual	0.0319	0.9047	0.9231