

Apriori Unimodality: Attributing polarization to sociodemographic groups

Dimitris Tsirmpas, John Pavlopoulos

March 2025

1 Introduction

Annotations are essential for a wide range of tasks, especially in fields such as content moderation, toxicity detection and hate speech detection. These tasks are especially essential for training systems that protect vulnerable and minority groups in online spaces, where they are often targeted [18, 4]. However, annotations for such tasks are subjective, and are usually based on majority-group annotators. This may limit the usefulness of the data, or even lead to biased and misconfigured systems.

The core of the problem is simple; if a comment targets marginalized groups, there can be cases where most annotators overlook it, while the few marginalized annotators vehemently disagree in vain. An example of such cases would be racist comments that are presented with coded language (usually referred to as “dog-whistles” [13]), which are not picked up by most annotators, but only by ones belonging in the targeted minority. Inversely, majority-group annotators may bias their own annotations against the minority-group annotators; CITATION for instance shows that toxicity models disproportionately marked African American speech as “toxic”. This issue exists even in representative samples.

Inter-annotator agreement is often used to detect such cases [8]. Metrics based on agreement mostly measure annotation variance, which may not be an optimal problem formulation for detecting such cases. *Polarization* is a better instrument in this regard, since it formulates annotation analysis as a cluster identification problem [11, 12]—specifically whether two or more annotation clusters are present. Figure 1




Figure 1: Disagreement is similar to variance; thus it overlooks cases where annotators may be “split down the middle”. Polarization on the other hand, is able to identify multiple clusters, even if they are close together.

displays how the two methods handle four common cases of annotation distributions for a hypothetical toxicity annotation task.

While we can detect polarization [12, 11], there is little work on how to detect whether it is actually caused by marginalized groups. Pavlopoulos and Likas [12] propose a mechanism which only works in the extreme case where overall polarization is zero; in this work we demonstrate that this case is not fre-

quent in real-world data. We instead propose the “*Apunim Metric*”, which attributes polarization to individual annotator sociodemographic groups. We provide a formulation which avoids inherent statistical limitations. Our contributions are:

1. We introduce a new metric which attributes polarization to specific sociodemographic groups.
2. We discover issues not yet acknowledged in literature, such as the presence of inherent polarization and conflicting polarization directions. Our proposed metric is designed to avoid these issues.
3. We apply our test to four datasets; two with human comments and annotations [15, 9] as well as two generated and annotated by Large Language Model (LLM) agents [17].

2 Methodology

In this section, we provide a formal mathematical formulation for the problem of attributing polarization to specific annotator characteristics (§2.1), and offer an intuitive rationale for how established polarization metrics can be leveraged (§2.2). We then introduce a data-point-level statistic that attributes polarization to individual SocioDemographic Background (SDB) groups (§2.3), and subsequently develop a statistic that generalizes this mechanism to an entire dataset.

2.1 Problem Formulation

SocioDemographic Backgrounds Since our goal is to pinpoint which specific characteristics contribute to polarization, we need to isolate individual groups within a SDB. Let Θ be one of multiple “dimensions” the set of all annotator SDB groups for a single dimension (e.g. “male”, “female” and “non-binary” if we are investigating annotator gender).

Annotations Now let $d = \{c_1, c_2, \dots\}$ be a dataset composed of multiple annotated data points. We assume that annotating a data point depends on two variables: inherent (“apriori”) controversy, and

the (“aposteriori”) annotator’s SDB, as well as uncontrolled factors such as mood and personal experiences (noise). Assuming that each data-point c is assigned multiple annotators, we can define the annotation multi-set $A(c) = \{(a, \theta)\}$, where a is a single annotation. As an example, the annotations for a comment in a sentiment analysis task would be formulated as:

$$A(\text{"Could be better, could be worse"}) = \{(\text{positive}, \text{female}), (\text{negative}, \text{male}), (\text{positive}, \text{female}), \dots\} \quad (1)$$

Polarization Each set of annotations features a certain degree of *polarization*. Unpolarized annotations are usually unimodal; the annotators disagree on the details (which would be shown roughly as a bell curve around the median annotation), not fundamentally (which would likely be shown as a multimodal distribution). Pavlopoulos and Likas [12] create a polarization metric, the “normalized Distance From Unimodality (nDFU)”, which measures whether an annotation set shows no polarization (unimodal distribution — $nDFU \rightarrow 0$), up to complete polarization (multimodal distribution — $nDFU \rightarrow 1$). Thus, we need a metric that measures how much the nDFU of a data point can be partially explained by θ (given that polarization exists).

2.2 Quantifying changes in polarization

Single data point polarization Intuitively, θ partially explains the polarization of a data point c when the annotations grouped by θ are more polarized compared to the full set of annotations. Figure 3 exhibits a hypothetical example where a misogynistic comment is annotated for toxicity by male and female annotators.¹ The annotations are generally polarized ($nDFU_{all} = 0.625$). The annotations by female annotators exhibit low polarization between themselves ($nDFU_{women} = 0.1$), since most agree the data-point

¹The use of only two predominant genders is made only for the purposes of demonstration.

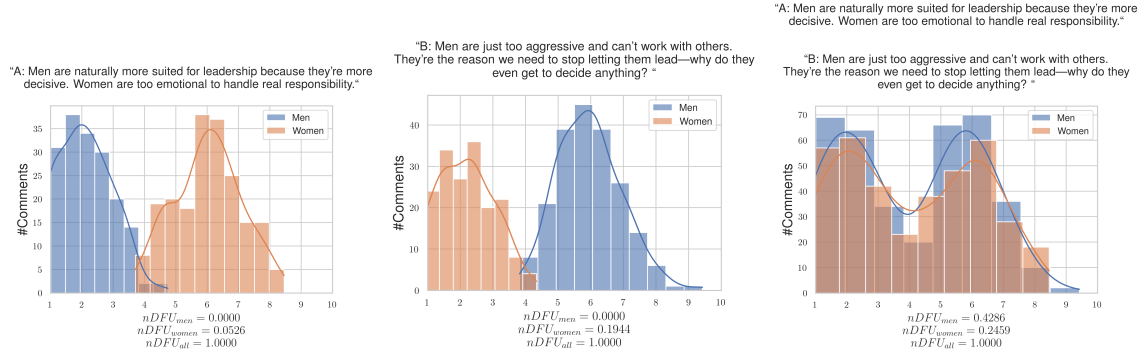


Figure 2: Hypothetical example of a polarizing discussion with two comments, for which the annotators disagree on which is the toxic one. If we aggregate the two comments, the polarization scores for both men and women significantly rise, obscuring whether the exhibited polarization can be partially attributed to gender.




Figure 3: Hypothetical example of a polarizing comment, where male and female annotators agree between themselves, but disagree with the opposite gender. Overall polarization is much greater than the polarization exhibited by the annotations grouped by gender.

is toxic. The set of male annotations, also shows low polarization ($nDFU_{men} = 0.375$), but for the opposite reason—most men agree that it is *not* toxic. This suggests that the overall polarization is driven by disagreements between male and female annotators.

Dataset polarization Given this observation, we would be tempted to aggregate all annotations in the dataset and compare the polarization of each θ with the full set of annotations $A(c)$. In fact this is very tempting since aggregating annotations solves the very real and important problem of annotation scarcity; each data point may be annotated by only 5 annotators, leading to a best-case 3-2 split between groups.

However, this formulation would not work well. Figure 2 illustrates a hypothetical discussion with two comments, both of which are toxic, but where male and female annotators disagree on *which* comment is the toxic one. If we aggregate the annotations to a single discussion, the opposing polarization effects might balance each other out, leading to a false negative. This example demonstrates that polarization has a *direction*, which is not captured by the $nDFU$ metric which is symmetric. Thus, care must be taken to not combine polarization effects of opposite directions.

2.3 The pol-statistic

As demonstrated above, it is necessary to define a metric for the polarization attribution of each individual data point. We thus introduce the “*pol*-

statistic”, which quantifies the change in nDFU when annotations are grouped by $\theta \in \Theta$.

In order to define this statistic, we need to measure (1) the polarization of a data-point when grouped by θ , and (2) the “apriori” polarization § 1 present in the annotations, which is driven by either the inherent controversy of the data-point. Note that we can not directly compare the θ -subset of annotations with the full set of annotations, since any statistical comparison between them violates the assumption of independent samples.

The observed polarization of a data-point on the group of annotations θ can be directly computed as:

$$pol(c, \theta) = nDFU(P(A(c), \theta)) \quad (2)$$

where $P(A(c), \theta_i) = \{a(c; \theta) \in A(c) | \theta = \theta_i\}$ is the partition of A for a data-point c , for which its annotations belong to the SDB group θ_i .

We can estimate the apriori polarization in data point c by bootstrapping. We randomly partition the annotations in $|\Theta|$ groups, with matching group sizes (e.g., if we have 100 annotations, 80 of which are made by male annotators and 20 by female annotators, we will create random partitions of sizes 80 and 20). The randomly partitioned annotations are then sampled t times. Formally, the expected inherent polarization will be given by:

$$\frac{1}{t} \sum_{i=1}^t nDFU(\tilde{P}_i(A(c))) \quad (3)$$

where $\tilde{P}_i$ is the random partition operator with matching group sizes.

2.4 Aposteriori Unimodality

By obtaining the pol-statistics for all data-points in a dataset d for SDB θ , we can get the mean aposteriori polarization for the dataset:

$$pol_{apost}(d, \theta) = \frac{1}{|d|} \sum_{c \in d} pol(c, \theta) \quad (4)$$

and similarly the mean apriori dataset polarization:

$$pol_{apriori}(d) = \frac{1}{|d|} \sum_{c \in d} \frac{1}{t} \sum_{i=1}^t nDFU(\tilde{P}_i(A(c))) \quad (5)$$




Figure 4: An overview of the Aposteriori Unimodality Test. We gather statistical information from the annotations (different SDBs are denoted by different colors) via the nDFU measure, which is aggregated on the data-point-level by the pol-statistic. The Aposteriori Unimodality Test is applied on the discussion level, operating on the pol-statistics of the individual data-points.

We can then define the *Apunim Statistic* as:

$$apunim(d, \theta) = \frac{pol_{apost}(d, \theta) - pol_{apriori}(d)}{1 - pol_{apriori}(d)} \quad (6)$$

If $apunim(d, \theta) \approx 0$, the exhibited polarization among annotators of group θ does not surpass what would be expected by chance. If $apunim(d, \theta) > 0$, then the group partially explains a rise in polarization, and inversely for the case $apunim(d, \theta) < 0$. The scope of the statistic and its relationship with the previously presented statistics are demonstrated in Figure 4.

Like most metrics, we also include a p-value alongside the κ value. This can be computed either non-parametrically, by repeatedly shuffling annotations and computing a null distribution of $apunim$, or parametrically, by assuming a normal distribution (which is generally a safe assumption if $t \geq 30$). In the latter case, we compute:

$$z = \frac{\hat{\kappa} - \mathbb{E}[\kappa]}{SE(\hat{\kappa})}, p = 1 - \Phi(z) \quad (7)$$

2.5 Technical Details

Since we are simultaneously considering $|\Theta|$ hypotheses, we apply a multiple comparison correction to the resulting p-values. We choose the Holms method [6]. The test is parameterized by the Family-Wise Error Rate (FWER), which is used to tune the strength of the correction; we can increase this value to make our test more conservative towards multiple hypotheses [3]). In general, it is safe to set FWER equal to the significance level of our test (e.g., $FWER = 0.95$ if we are looking for $p < 0.05$).

3 Results

We apply our test on four datasets; two human-annotated datasets, and two with comments and annotations generated by LLMs.

3.1 Human Datasets

We use the datasets provided by Kumar et al. [9] and Sap et al. [15]. The datasets feature various online comments, each annotated by a number of human annotators (5 and 4 – 6 annotators per comment respectively). The Sap et al. [15] dataset includes racism annotations for 626 Twitter/X comments, and standard SDB information (age, race, education, gender). The Kumar et al. [9] dataset measures toxicity on various online comments, and the provided SDBs include mostly annotator experiences (e.g., whether they have been personally targeted online) as well as standard SDB information such as sexual orientation, age and education. Since the dataset is extensive ($XXXX$ tweets), we only select a sample of 5000 comments. This is due both to computational constraints, as well as statistical tests being unreliable on large enough samples [16]. Assuming a confidence level of $\alpha = 0.05$ we test whether any of the provided SDB dimensions can partially explain polarization in the toxicity/racism annotations.

The results for the Kumar et al. [9] and Sap et al. [15] datasets can be found in Tables ??, ?? respectively.

Table 1: Aposteriori Unimodality kappa and pvalue results for the Sap et al. 2022 dataset

SDB Feature		apunim	pvalue_parametric	pva
Age	(-0.001, 20.0]	0.000000	NaN	
	(20.0, 40.0]	-0.028003	0.946350	
	(40.0, 60.0]	0.035342	0.934080	
	(60.0, 80.0]	0.078341	0.757341	
Ethnicity	black	-0.419412	0.413176	
	hisp	0.000000	NaN	
	other	0.000000	NaN	
	white	0.149386	0.690118	
Gender	man	0.151317	0.711524	
	nonBinary	0.000000	NaN	
	woman	-0.349321	0.413578	

3.2 Synthetic Datasets

We use two synthetic datasets; one is the Virtual Moderation Dataset (VMD) presented in Tsirmpas, Androutsopoulos, and Pavlopoulos [17]. This dataset features 140 discussions, each having 14 comments (and usually up to 14 facilitator comments), each comment annotated by 10 LLM annotators, each supplied with a different SDB. The second dataset, is an extension of VMD, where we select four random unmoderated discussions, and employ 100 distinct LLM annotators.

We find no statistically significant results in any of the SDB dimensions, in any of the datasets (annotator age, gender, sexual orientation, education, employment and political alignment). While polarization does exist in the dataset, it can not be attributed to any of the synthetic SDB groups. This suggests that SDB prompting can not be used to simulate annotations for human groups, which is consistent with relevant literature on the topic [1, 5, 14, 7, 2, 10].

The results for the VMD and 100-annotator datasets can be found in Tables ??, ?? respectively.

3.3 Effect of number of annotators

In §2.4 we mentioned that the test relies on enough annotations for each SDB group and data-point. This

Table 2: Aposteriori Unimodality kappa and pvalue results for the Kumar et al. 2021 dataset

SDB Feature		apunim	pvalue_parametric	pvalue_nonparametric
Age	18 - 24	-0.030410	0.954856	0.967112
	25 - 34	0.040010	0.948614	0.963320
	35 - 44	-0.010025	0.987850	0.998125
	45 - 54	0.018776	0.973235	0.971674
	55 - 64	0.003176	0.994370	1.000000
	65 or older	0.017648	0.955945	1.000000
	Prefer not to say	0.000000	NaN	apunim0.000000pvalue.p
Education	Associate degree	0.005727	0.923705	0.947000
	Bachelor's degree	0.095408	(14.927, 33.25)	0.908507
	College, no degree	-0.108984	(33.25, 51.5)	0.927355
	Doctoral degree	0.008115	(51.5, 69.75)	1.000000
	High School graduate	-0.040907	(69.75, 88.0)	0.021639
	Master's degree	0.054708	high-school	0.980510
	No high school	-0.007941	none	0.034102
	Other	0.000000	university	1.000000
	Prefer not to say	0.000000	blue-collar	-0.017141
	Professional degree	0.001086	unemployed	1.000000
Has Been Targeted	False	-0.095286	white-collar	-0.002428
	True	-0.112036	asian	-0.025802
Is Parent	No	-0.136264	black	0.840822
	Prefer not to say	-0.004510	other	-0.035818
	Yes	0.093870	white	0.834386
Is Transgender	No	-0.151712	female	0.068891
	Prefer not to say	0.515608	male	0.852657
	Yes	0.016393	non-binary	-0.060608
Political Affiliation	Conservative	0.066361	apolitical	0.776316
	Independent	-0.016543	left-wing liberal	0.002739
	Liberal	-0.039058	right-wing conservative	0.416667
	Other	-0.018498	bisexual	0.969231
	Prefer not to say	-0.001853	homosexual	-0.107289
Seen Toxicity	False	0.145958	other	0.974428
	True	-0.117346	straight	0.992056
Sexual Orientation	Bisexual	-0.179856	0.786152	0.977920
	Heterosexual	-0.131972	0.808635	1.000000
	Homosexual	-0.044238	0.940537	0.074502
	Other	0.003322	0.979759	0.000324
	Prefer not to say	-0.028099	0.957889	0.810560
	NaN	0.077491	0.756128	0.831252
				-0.011836
Thinks Religion Is Important	Not important	-0.167580	0.790175	0.858974
	Not too important	-0.026466	0.961419	0.982239
	Prefer not to say	0.004367	0.983731	1.000000
	Somewhat important	0.036322	0.954503	0.975619
	Very important	0.088742	0.893641	0.964399
Thinks Toxicity Is Problem	Frequently a problem	-0.057269	0.927206	0.983830
	Not a problem	0.039260	0.934240	0.997835
	Occasionally a problem	-0.009915	0.987592	0.996255
	Rarely a problem	0.044619	0.943330	0.956728
	Very frequently a problem	-0.057531	0.914521	0.954014

Table 3: Aposteriori Unimodality kappa and pvalue results for the 100 Annotator Synthetic dataset




Figure 5: Standard error of the *pol-statistic* when sampled with replacement from the 100-annotator synthetic dataset for various number of annotators.

can be an issue, given the cost required to utilize multiple human annotators for each data-point [14]. Figure 5 demonstrates the effect of the number of annotators to the *pol-statistic* estimation. We use the 100-annotator synthetic dataset and sample progressively 3 – 100 annotators with replacement, then calculate the standard error with the mean *pol-statistic*. As expected, the standard error is inversely proportional to the number of annotators, although the difference is not great, even for a low number of annotators.

4 Conclusion

We introduced the “Aposteriori Unimodality Test”, a statistical test that detects whether certain sociodemographic groups partly cause polarization in annotations. Our test is resistant to low samples of annotations, bypasses common statistical issues such as sample dependence, and avoids newly-discovered issues such as the presence of inherent polarization and polarization direction. We apply our test to two human-generated and two LLM-generated datasets and find interesting patterns on the former, and verify current literature on LLM SDB prompting on the latter.

5 Limitations

Since there are no other established tests that attribute polarization to sociodemographic characteristics, there is no way to verify that the results presented in §3 are accurate. Similarly, there is no qualitative way of evaluating our test, other than observations on the (non-) efficacy of LLM SDB prompting and our own intuition. Lastly, while our test seems to work sufficiently well with a small number of annotators per data-point, we still encourage researchers using this test to have at least a few annotators of each sociodemographic group that is being tested, in their dataset.

6 Ethical Considerations

While our work aims to help protect marginalized and disadvantaged groups by attributing polarization to certain subgroups, it can be taken advantage of by malicious actors. Given a dataset with fine-grained SDB information, these actors can use our test to target specific vulnerable groups. We thus urge researchers to keep datasets including such information protected, and provided to others only under explicit permission.

References

- [1] Jacy Reese Anthis et al. *LLM Social Simulations Are a Promising Research Method*. 2025. arXiv: 2504.02234 [cs.HC]. URL: <https://arxiv.org/abs/2504.02234>.
- [2] James Bisbee et al. “Synthetic Replacements for Human Survey Data? The Perils of Large Language Models”. In: *Political Analysis* 32.4 (2024), 401–416. DOI: 10.1017/pan.2024.5.
- [3] S.Y. Chen, Z. Feng, and X. Yi. “A general introduction to adjustment for multiple comparisons”. In: *Journal of Thoracic Disease* 9.6 (2017). doi: 10.21037/jtd.2017.05.34; PMID: 28740688; PMCID: PMC5506159, pp. 1725–1729.

- [4] Shelby Dioum. *What Explains the Increase in Online Hate Speech?* Accessed: 2025-09-22. University of California, Davis, Mar. 2025. URL: <https://www.ucdavis.edu/magazine/what-explains-increase-online-hate-speech>.
- [5] Luke Hewitt et al. “Predicting Results of Social Science Experiments Using Large Language Models”. Equal contribution, order randomized. Aug. 2024.
- [6] Sture Holm. “A Simple Sequentially Rejective Multiple Test Procedure”. In: *Scandinavian Journal of Statistics* 6.2 (1979), pp. 65–70. ISSN: 03036898, 14679469. URL: <http://www.jstor.org/stable/4615733> (visited on 07/07/2025).
- [7] Bernard J. Jansen, Soon gyo Jung, and Joni Salminen. “Employing large language models in survey research”. In: *Natural Language Processing Journal* 4 (2023), p. 100020. ISSN: 2949-7191. DOI: <https://doi.org/10.1016/j.nlp.2023.100020>. URL: <https://www.sciencedirect.com/science/article/pii/S2949719123000171>.
- [8] Petra Kralj Novak et al. “Handling Disagreement in Hate Speech Modelling”. In: *Information Processing and Management of Uncertainty in Knowledge-Based Systems*. Ed. by Davide Ciucci et al. Cham: Springer International Publishing, 2022, pp. 681–695. ISBN: 978-3-031-08974-9.
- [9] Deepak Kumar et al. “Designing toxic content classification for a diversity of perspectives”. In: *Proceedings of the Seventeenth USENIX Conference on Usable Privacy and Security*. SOUPS’21. USA: USENIX Association, 2021. ISBN: 978-1-939133-25-0.
- [10] Terrence Neumann, Maria De-Arteaga, and Sina Fazelpour. *Should you use LLMs to simulate opinions? Quality checks for early-stage deliberation*. 2025. arXiv: 2504.08954 [cs.CY]. URL: <https://arxiv.org/abs/2504.08954>.
- [11] John Pavlopoulos and Aristidis Likas. “Distance from Unimodality for the Assessment of Opinion Polarization”. In: *Cognitive Computation* 15.2 (2023), pp. 731–738. ISSN: 1866-9964. DOI: 10.1007/s12559-022-10088-2. URL: <https://doi.org/10.1007/s12559-022-10088-2>.
- [12] John Pavlopoulos and Aristidis Likas. “Polarized Opinion Detection Improves the Detection of Toxic Language”. In: *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Yvette Graham and Matthew Purver. St. Julian’s, Malta: Association for Computational Linguistics, Mar. 2024, pp. 1946–1958. URL: <https://aclanthology.org/2024.eacl-long.117/>.
- [13] Anne Quaranto. “Dog whistles, covertly coded speech, and the practices that enable them”. In: *Synthese* 200.4 (2022), p. 330.
- [14] Luca Rossi, Katherine Harrison, and Irina Shklovski. “The Problems of LLM-generated Data in Social Science Research”. In: *Sociologica* 18.2 (2024), 145–168. DOI: 10.6092/issn.1971-8853/19576. URL: <https://sociologica.unibo.it/article/view/19576>.
- [15] Maarten Sap et al. “Annotators with Attitudes: How Annotator Beliefs And Identities Bias Toxic Language Detection”. In: *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Ed. by Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz. Seattle, United States: Association for Computational Linguistics, July 2022, pp. 5884–5906. DOI: 10.18653/v1/2022.naacl-main.431. URL: <https://aclanthology.org/2022.naacl-main.431/>.
- [16] David Trafimow et al. “Manipulating the Alpha Level Cannot Cure Significance Testing”. In: *Frontiers in Psychology* 9 (2018), p. 699. DOI: 10.3389/fpsyg.2018.00699.

- [17] Dimitris Tsirmpas, Ion Androutsopoulos, and John Pavlopoulos. *Scalable Evaluation of On-line Facilitation Strategies via Synthetic Simulation of Discussions*. 2025. arXiv: 2503.16505 [cs.HC]. URL: <https://arxiv.org/abs/2503.16505>.
- [18] United Nations. *Targets of Hate*. Accessed: 2025-09-22. Sept. 2025. URL: <https://www.un.org/en/hate-speech/impact-and-prevention/targets-of-hate>.

A Appendix

A.1 Acronyms

VMD	Virtual Moderation Dataset
LLM	Large Language Model
SDB	SocioDemographic Background
nDFU	normalized Distance From Unimodality
FWER	Family-Wise Error Rate

A.2 Demonstrative examples

Our paper uses synthetic examples to demonstrate the intuition behind the Aposteriori Unimodality test. These examples use random sampling of manually selected distributions. Table 4 summarizes the synthetic annotation distributions used to demonstrate the Aposteriori Unimodality test. In each scenario, annotations for men and women were sampled independently from Gaussian distributions with distinct parameters.

Figure	Group	Distribution	Size
Figure 3	A_{men}	$\mathcal{N}(2, 1.3)$	50
	A_{women}	$\mathcal{N}(8, 1.3)$	50
Figure 2 (1st comment)	A_{men}	$\mathcal{N}(6, 1)$	200
	A_{women}	$\mathcal{N}(2, 1)$	200
Figure 2 (2nd comment)	A_{men}	$\mathcal{N}(2, 1)$	200
	A_{women}	$\mathcal{N}(6, 1)$	200

Table 4: Synthetic annotation sets used to illustrate the Aposteriori Unimodality test. Each group corresponds to samples drawn from a Gaussian distribution with specified mean and standard deviation.