

Detecting, quantifying, and attributing fundamental disagreements across annotator groups

Anonymous ACL submission

Abstract

Current annotation agreement metrics are often sensitive to group size imbalances and noise, making them inadequate for many subjective tasks such as toxicity and hate-speech detection, where minority opinions are most important. Polarization has recently been proposed as a more robust way to distinguish minor disagreements from systematic differences in opinion, but existing approaches do not provide practical tools for attributing polarization to specific annotator groups. We evaluate current methods and identify two major limitations in realistic settings: (1) the presence of “inherent” polarization that cannot be attributed to any known or latent groups, and (2) opposing polarization effects canceling each other out in aggregated annotations. To address these issues, we introduce a new metric that measures and tests the statistical significance of polarization attribution for annotator groups while avoiding these limitations. We apply our method to four subjective NLP datasets and find that gender and race consistently explain polarization patterns, while differences between annotator groups become stronger as the groups are further apart. Finally, we estimate the minimum number of annotators required for reliable results and release an open-source Python library implementing our approach.

1 Introduction

Annotations are essential for a wide range of tasks, especially in fields with subjective judgements, such as content moderation, toxicity, and hate speech detection. These tasks are especially important for training supervised learning models that remove or flag comments targeting vulnerable and minority groups in online spaces, where they are often vulnerable (United Nations, 2025; Dioum, 2025). However, annotations for such tasks are subjective and typically aggregated based on majority-voting (Ivey et al., 2025; Giannoutsos et al., 2025; Kralj Novak et al., 2022; van der




Figure 1: Simulated experiment exploring four scenarios of a hate speech detection task applied to a controversial statement by D. Trump (Jan. 2019). The variance in annotations across different annotators—Muslim annotators in the upper plots—may cause traditional agreement metrics may fail to capture distinct group biases (Pavlopoulos and Likas, 2024), especially when mixed opinions occur (top-right).

Velden et al., 2025). This may at best limit the usefulness of the data, or at worst lead to biased and misconfigured systems (Sap et al., 2019; van der Velden et al., 2025).

The core of the problem is simple; if a comment targets marginalized groups, there can be cases where most annotators overlook it, while the few marginalized annotators vehemently disagree with the majority. These *fundamental* disagreements are typically ignored via aggregation (e.g., majority label voting) or filtering (e.g., keeping only items with large overall annotator agreement). For example, coded language may not be picked up by most annotators but only by ones belonging in the targeted minority (Quaranto, 2022), as shown in Fig. 1. Furthermore, majority-group annotators may themselves be biased against minority groups, e.g., toxicity models disproportionately marking African American (AA) speech as “toxic” (Sap et al., 2019).

Instead of using inter-annotator agreement, we could use *polarization*, which formulates anno-

tation analysis as a cluster identification problem (Pavlopoulos and Likas, 2023, 2024). While we can detect polarization (Pavlopoulos and Likas, 2024), there is little work on how to detect whether it is *actually caused by marginalized groups*. Our work thus tries to answer the following Research Question:

RQ. *Can we attribute polarization to annotator subgroups?*

We first identify limitations of current work on polarization:

1. Existing mechanisms (Pavlopoulos and Likas, 2024) operate only as binary classifiers, producing non-quantifiable/comparable results.
2. These mechanisms only test for the extreme case where *all* present polarization can be attributed to a single annotator group, which we find to be extremely rare in real-life datasets (App. B).
3. We are the first to report the existence of *inherent polarization*: most comments having some degree of polarization that cannot be attributed to *any possible partition* of annotators (Finding 1).
4. We also find that annotator groups may exhibit opposing polarization attribution between comments, requiring isolating annotations to each individual comment when analyzing whole datasets (Finding 2).

We then introduce a new metric (“Aposteriori Unimodality–*apunim*”) that attributes polarization to specific annotator groups without and develop a parametric test for statistical significance. Our metric is generalizable across any domain and modality (e.g., video/image classification) and, unlike previous approaches, enables *quantitative* comparisons across datasets. By considering the entire set of comments in a dataset *apunim* helps to identify genuine polarization and distinguish it from chance attribution, which is essential when analyzing large datasets.

Having established our metric, we apply it to four Natural Language Processing (NLP) datasets focused on Toxicity and Hate Speech detection (§5), finding that annotator gender and race significantly impacts annotation (Finding 3), and that fundamental disagreements grow more pronounced the further apart annotator groups are (e.g., irreligious vs. very religious annotators–Finding 4).

Finally, we release an open source Python library that implements our metric, which can be directly installed via PyPi.¹ and provide an estimation on the minimum number of annotators required for robust estimates (§6). The code for the experi-

ments, graph creation, and analysis can be found on our project’s repository.²

2 Related Work

2.1 Agreement

There are many popular metrics for measuring inter-annotator agreement such as Cohen’s kappa, Krippendorff’s alpha and Fleiss’s kappa. All of these metrics attempt to distinguish genuine disagreement among annotators from noise (or “random disagreement”) using statistical methods. However, most cannot be generally used for comparing intergroup annotation agreement without extending them (e.g., computing pairwise Cohen’s kappa for all groups), or placing constraints on when they can be used (e.g., no missing labels and same number of annotations per annotator for Fleiss’s kappa). They have also been criticized for ignoring fundamental differences that can arise from different social, cultural and demographic backgrounds (Feinstein and Cicchetti, 1990; Byrt et al., 1993; van der Velden et al., 2025; Checco et al., 2017), for their inability to quantify and compare results across experiments (Wong et al., 2021), and for fundamentally faulty assumptions around the concept of “chance-correction” (Checco et al., 2017).

Recent work has attempted to work around these limitations by using simpler statistical models (van der Velden et al., 2025), adapting Cohen’s kappa for subgroup agreement (Wong et al., 2021), or creating a Bayesian model to explicitly differentiate between subgroup disagreement and noise (Ivey et al., 2025). While these extensions do improve the performance of these metrics, they are typically less explainable and intuitive.

2.2 Polarization

There are multiple competing formulations for polarization. One of these approaches (Akhtar et al., 2019) utilizes χ^2 to estimate in-group vs out-group variance, but still necessitates traditional agreement metrics for analysis. Others attempt to solve this problem by framing disagreement as a cluster identification problem, applying either parametric (Checco et al., 2017), or non-parametric (Pavlopoulos and Likas, 2024; Mignemi et al., 2024) approaches to the annotation histograms.

For our work, we use a non-parametric polarization metric, normalized Distance From Unimodality (nDFU), which has been shown to correlate better with human judgment of disagreement than established metrics (Pavlopoulos and Likas, 2024). nDFU takes values in the range $[0, 1]$, where values close to 0 indicate little or no polarization (i.e.,

¹Library: <https://anonymous.4open.science/r/apunim-5F70/README.md>
Documentation: <https://anonymous.4open.science/r/apunim-page/>

²Replication code: <https://anonymous.4open.science/r/aposteriori-unimodality-E8C1/>

opinions are concentrated around a single mode), while values close to 1 indicate strong polarization with multiple distinct peaks. The metric is defined as follows:

$$\text{nDFU}(X) = \frac{\max_{i \neq m} \left(f_i - \begin{cases} f_{i-1}, & \text{if } i > m \\ f_{i+1}, & \text{if } i < m \end{cases} \right)}{f_m} \quad (1)$$

$$m = \arg \max_i f_i$$

where, X denotes the set of annotations, f_i the relative frequency of the i -th annotation, and m the histogram peak. Intuitively, nDFU measures how prominent secondary peaks are compared to the main peak.

2.3 Attributing polarization to groups

While many studies study the effect of subgroups on annotation agreement (Goyal et al., 2022; Binns et al., 2017; Giorgi et al., 2025; Al Kuwatly et al., 2020; Sachdeva et al., 2022), attributing polarization to subgroups is a relatively new idea. Under the name *aposteriori unimodality* (Pavlopoulos and Likas, 2024), the original idea is that if polarization in a single item is positive, but zero for all annotator subgroups, then those subgroups are “causing” the exhibited polarization. Though promising, this mechanism operates on an edge-case scenario where *all* polarization is explained by a single split. Additionally, getting polarization estimates from few annotations per item is extremely noisy, which is especially problematic when the output is binary (i.e., an item either is or is not polarizing) instead of a quantifiable value.³

3 What are the limitations of polarization attribution?

3.1 Problem Formulation

Let $d = \{c_1, c_2, \dots\}$ be a dataset composed of multiple annotated comments.⁴ Each comment c is assigned multiple annotators, which creates the annotation multi-set $A(c) = \{(a_1, g_1), (a_2, g_2), \dots\}$, where a is a single annotation. As an example, the annotations for a comment in a sentiment analysis task where we group the annotators by gender, would be formulated as:

$$A(\text{“Could be better, could be worse”}) \\ = \{ (+, F), (-, M), (+, F), \dots \}$$

³For example, when presented with noisy estimates of a metric in the $[0, 1]$ range, we would ideally want to distinguish between cases where attribution is close to 0.001, and cases where attribution is 0.49.

⁴These items could be of any modality, such as comments in a discussion, images, videos or survey questions, even though we consider comments in this study.




Figure 2: Synthetic experiment simulating a polarizing comment, where male and female annotators agree between themselves, but disagree with the opposite gender.

Intuitively, a group partially explains the polarization of a comment c when the annotations of that group $A(c | g) = \{a(c; g') \in A(c) | g' = g\}$ exhibit higher polarization compared to the full set of annotations $A(c)$. Consider the (simulated) case where a misogynistic comment is annotated for toxicity by male and female annotators shown in Fig. 2. The annotations are generally polarized ($\text{nDFU}_{all} = 0.62$); but both female ($\text{nDFU}_{women} = 0.10$) and male ($\text{nDFU}_{men} = 0.37$) annotators exhibit low polarization between their respective groups. This pattern suggests that the overall polarization is driven by substantive disagreements between male and female annotators.

3.2 Inherent polarization




Figure 3: Apriori (inherent rather than observed) polarization across our datasets. Contrary to assumptions made in prior work, inherent polarization is rarely zero and can even reach $\text{nDFU} = 1$.

Inherent polarization refers to disagreement that

cannot be attributed to *any specific* grouping of annotators; that is, polarization that exists no matter how we split the annotators in any arbitrary group. Prior work (Pavlopoulos and Likas, 2024) attributed polarization to annotator groups only when the observed polarization for each individual group became exactly zero (e.g. $nDFU(all) > nDFU(men) = nDFU(women) = 0$). The presence of inherent polarization would invalidate this theory, since there would be no possible group g s.t. $nDFU(g) = 0$.

Definition 1: Inherent polarization. *Polarization that cannot be explained by any annotator known or latent annotator grouping.*

Formally, the inherent polarization for a single comment c is given by:

$$Pol_{inh}(c) = \min_i \{nDFU(\tilde{r}_i(A(c)))\} \quad (2)$$

where $\tilde{r}$ is the random partition operator.

To examine this assumption, we create random annotator groupings for each comment in our datasets. Specifically, for every comment, we randomly partition the annotators into random groups (with randomly selected sizes) 100 times and compute the observed polarization for each partition. We then take the *minimum* observed polarization scores across all random partitions, which guarantees that we check for all possible variables, known and latent. Repeating this process for every comment produces the distribution shown in Fig. 3, demonstrating that inherent polarization is consistently non-zero and can even be substantial for many comments.

Finding 1. *Each comment may exhibit a degree of inherent polarization that cannot be attributed to any known or latent annotator grouping.*

3.3 Conflicting polarization

Previous work (Pavlopoulos and Likas, 2024) operated on the comment level. In practice however, this approach is inherently noisy due to a lack of information (i.e., limited annotations) relating to each individual comment. Many datasets in NLP are annotated by only 3–5 annotators per comment, which is especially problematic when we want to analyze patterns in annotator subgroups; in a best-case scenario of 5 annotators split evenly between two groups, we would get a 3-2 split. This issue is especially prevalent for polarization, since $nDFU$ is undefined for groups of two (and arguably marginally useful for groups of three), since it requires at least two histogram peaks to function (see Eq. 1), a limitation not acknowledged in the original work.

This limitation leaves us with two choices: (1) only use datasets featuring at least 6–8 annotators, or (2) aggregate annotations across comments for

each dataset. The former case is unrealistic, especially for most NLP datasets due to annotation costs. We consider the latter case in an experiment shown in Fig. 4 where we combine annotations across two comments. We observe that, should the two comments be polarized for the opposite reason, the opposing polarization effects might balance each other out, leading to a false negative. Therefore, we should never aggregate annotations across comments.⁵

Finding 2. *Polarization has a direction, which may compromise cross-comment annotation aggregation.*

4 How do we fix it?

4.1 From inherent to apriori polarization

In §3.1, we explained how our methodology needs to compare the polarization of $A(c)$ and $A(c | g)$. Direct comparisons are not advisable, since we would be comparing statistics between a set and its subset—in this case, standard practice dictates comparing the subset with its complement; i.e., the in-group vs. the out-group instead of the in-group vs. all the annotations. However, this also would not work; if we apply the subset-vs-complement method in the example of Fig. 2, we would end up subtracting the $nDFUs$ of men and women, which are similar as they are both unimodal distributions, falsely concluding that there is no polarization present.

We instead compare the observed polarization of $A(c)$ with its inherent polarization, following the methodology introduced in §3.2 with two key distinctions. (1) Instead of using random splits, we create random stratified splits.⁶ (2) Instead of finding the minimum possible polarization exhibited by any known or latent group, we calculate the *expected* (average). Therefore, we do not test whether the group is polarized against the whole; nor whether the group explains *some polarization* (as would be the case if we used Eq. 2), but rather whether the group explains polarization *more than other groups*, which is a much stricter test.⁷ Specifically, we define a-priori unimodality for a comment c as:

$$Pol_{apr}(c) = \frac{1}{t} \sum_{i=1}^t nDFU(\tilde{r}_i(A(c))) \quad (3)$$

Definition 2: A-priori polarization. *The expected polarization of a comment’s annotations.*

⁵A similar assumption was followed in early parametric approaches, which assumed different parameters for each comment (Checco et al., 2017).

⁶I.e., if we have 100 annotations, 80 of which are made by male annotators and 20 by female annotators, we create random partitions of sizes 80 and 20.

⁷Indeed, by definition, all groups exhibit equal or larger polarization than the inherent pol. value.




Figure 4: A synthetic experiment simulating a polarizing discussion with two comments where annotators disagree about which one is toxic. When we aggregate the comments, we unintentionally change the comparison: instead of comparing two unimodal distributions (each comment viewed separately) with two bimodal distributions (their aggregated annotations), we end up comparing two bimodal distributions, obscuring the source of the polarization.

4.2 Expanding to the whole dataset

As discussed in §3.3, we must not aggregate the annotations themselves before applying the $nDFU$ metric. We can however estimate how much polarization is generally present in the dataset, by simply averaging the *a-priori* polarization of its comments. Intuitively, a high apriori polarization score means that much of the polarization in the dataset is likely not driven by any group.

$$Pol_{apr}^d = \frac{1}{|d|} \sum_{c \in d} Pol_{apr}(c) \quad (4)$$

Which we will compare with the average observed polarization of all comments:

$$Pol_{obs}^d(g) = \frac{1}{|d|} \sum_{c \in d} nDFU(A(c | g)) \quad (5)$$

We can now compare the polarization that is attributed to a single-dimensional annotator group (Equation 5) to a prior one (Equation 4). We define our normalized metric, which we call “*apunim*” (*Aposteriori Unimodality*), as:

$$apunim(d, g) = \frac{Pol_{obs}^d(g) - Pol_{apr}^d}{1 - Pol_{apr}^d} \quad (6)$$

Definition 3: Apunim. The degree to which an annotator group contributes to overall polarization in a set of comments.

Since $nDFU \rightarrow [0, 1] \Rightarrow apunim \rightarrow [-1, 1]$. The interpretation of this metric boils down to three scenarios:

- $apunim(\theta) \approx 0$: The examined group does not meaningfully disagree with the other annotators.

- $apunim(\theta) > 0$: The examined group has fundamental disagreements with the other annotators.
- $apunim(\theta) < 0$: Overall polarization *drops* when we add annotators from this group to this task. This occurs when intragroup polarization is higher than apriori polarization; indicating that annotators of that group are polarized between themselves.

4.3 Statistical significance

Since annotation histograms may be noisy, we additionally assess whether the observed *apunim* value could arise by chance. To assess statistical significance, we construct a null distribution of *apunim* values from the randomized partitions and compare the observed *apunim* against this distribution using a one-sample Student-*t* test. Thus, the null hypothesis states that the observed polarization is indistinguishable from polarization induced by random annotator assignments:

$$H_0 : apunim(d, g) \sim apunim_{rand}(d, g)$$

$$H_a : apunim(d, g) \not\sim apunim_{rand}(d, g)$$

where $apunim_{rand}(d, g)$ denotes the distribution of *apunim* values obtained from randomized annotator partitions. For a detailed explanation of the p-value computation and the assumptions underlying the test, refer to App. C.

5 Attributing real-life polarization to human groups

5.1 Datasets

We use four datasets from current NLP research that include sociodemographic information at the level of individual annotators; specif-

Dataset	#Com	#Ann	#Dims	Task
Kumar	106,035	5	10	Tox
Sap	626	5–6	3	HS
DICES-350	72,103	70–76	4	HS
DICES-990	43,050	123	4	HS

Table 1: Datasets used in this study. $\#Ann$ refers to the number of annotations per comment (see App. B), and $\#Dims$ to sociodemographic dimensions (e.g., age, gender). Tasks are either toxicity (Tox) or hate speech detection (HS).

ically, the “Kumar” (Kumar et al., 2021) and “Sap” (Sap et al., 2022) datasets, which are concerned with Toxicity and hate speech detection respectively, and the DICES-350 and DICES-990 datasets (Aroyo et al., 2023), which focus on LLM response safety.⁸ While the toxicity and hate speech detection tasks are not equivalent, they share substantial conceptual and operational overlap, as both require annotators to assess the presence and severity of harmful or abusive language. Thus, we expect the sources of annotator disagreement and bias that arise in these tasks will be broadly similar.

A brief overview of dataset characteristics is provided in Table 1. The Kumar dataset stands out for its large size, featuring more than 100,000 annotated comments, as well as 10 distinct sociodemographic dimensions. In contrast, the DICES datasets contain fewer annotated comments but involve more than an order of magnitude more annotators per comment than the other datasets. The Kumar and Sap datasets contain a large number of polarizing items, unlike the DICES datasets, where most items are unpolarized (Fig. 7).⁹

5.2 Experimental Setup

Assuming a confidence level of $\alpha = 0.05$ we test whether any of the provided sociodemographic dimensions can partially explain polarization in the toxicity / hate speech annotations, as shown in Table 1.¹⁰ Note that some groups are not included, either because the results were not statistically significant, or because there were no polarized items annotated by more than one annotator of these groups. See App. E for details on other parameters and execution times, and App. F for the full results.

⁸For the latter, we consider only labels relating to racism (Q3_bias_overall), thereby transforming the task to hate speech detection.

⁹This systematic discrepancy could be attributed to the DICES datasets being composed of Large Language Model (LLM)-generated dialogues with already-aligned models tuned to avoid racist speech.

¹⁰We sample 10,000 comments from the Kumar dataset due to computational costs arising from the extraordinary amount of **pc**ls (**pc**ls).




Figure 5: Polarization attribution trends in ordinal groups. The left side of the x-axis (low orders) refers to young annotators, low education, low religiousness, and indifference towards the impacts of toxicity and technology. The full ordinal modes for each dimension and their individual labels can be found in App. F.

5.3 Which factors contribute to polarization?

Race Across all datasets apart from Sap, we observe a consistent pattern in which annotation polarization is linked to annotator characteristics, particularly ethnicity. In the hate speech detection datasets (DICES-350, DICES-990) *ethnicity is the strongest contributor to polarization*, with statistically significant effects for at least one major ethnic group in each dataset (AAs: -0.0428 /Asians: 0.0659 for DICES-350, Whites: -0.0377 for Kumar, and all ethnicities for DICES-990). These results align with prior work suggesting that ethnic background is a systematic source of annotator disagreement rather than isolated noise (Goyal et al., 2022), validating our initial motivation.

Gender Polarization is also consistently linked to gender across all datasets except DICES-350. In DICES-990, men and women show statistically significant effects in opposite directions (Men: 0.028 , Women: -0.025), mirrored in Sap, where men have a positive attribution score (0.091) and women a strongly negative one (-0.244). By comparison, the Kumar dataset shows almost identical attribution scores for female and male annotators (Female: 0.586 , Male: 0.588), a pattern that also holds for many other attributes in that dataset. This may indicate that *apunim* is partly task-dependent, or that the test becomes less discriminative when many comments are annotated by relatively few annotators. Nevertheless, these patterns mirror prior findings with respect to hate speech (Wojatzki et al., 2018) and toxicity (Binns et al., 2017) annotation.

Finding 3. *Annotator age and race are consistent causes of polarization.*

Group	apunim	support
DICES-350		
Race=AfricanAmerican	-0.043	9077
Race=Asian	0.066	8138
DICES-990		
Age=GenX+	-0.046	14384
Age=GenZ	0.063	13741
Edu=High school and below	-0.075	7419
Gender=Man	0.028	32494
Gender=Woman	-0.025	33471
Race=AfricanAmerican	0.091	5810
Race=Latino	-0.133	6610
Race=Other	0.075	24321
Race=White	-0.120	11392
Sap		
Gender=Man	0.091	1439
Gender=Woman	-0.244	877
Kumar		
Age=18–24	0.611	331
Age=25–34	0.606	8506
Age=35–44	0.557	2556
Age=45–54	0.555	441
Age=55–64	0.637	110
Edu=Masters	0.642	858
Edu=Bachelors	0.594	9013
Edu=Associate degree	0.585	335
Edu=College, no degree	0.570	1425
Edu=High School graduate	0.446	144
Eth=Asian	0.461	49
Eth=Black	0.618	468
Eth=White	0.546	26239
Gender=Female	0.586	14479
Gender=Male	0.588	11716
Past Targeted=No	0.551	25773
Past Targeted=Yes	0.608	3893
Parent=No	0.566	10784
Parent=Yes	0.593	15929
Trans=No	0.484	40404
Politics=Conservative	0.620	3075
Politics=Independent	0.574	3270
Politics=Liberal	0.580	7897
ReligionImportant=No	0.555	4041
ReligionImportant=Low	0.572	315
ReligionImportant=Mid	0.624	2542
ReligionImportant=High	0.631	5403
SeenToxicity=No	0.627	2454
SeenToxicity=Yes	0.529	28821
SexOrientation=Bi	0.621	467
SexOrientation=Het	0.528	32217
TechImportant=Negative	0.492	217
TechImportant=Neutral	0.587	426
TechImportant=Positive	0.579	13685
TechImportant=Very pos.	0.622	2942
ToxProblem=Never	0.451	55
ToxProblem=Rare	0.572	1007
ToxProblem=Occasionally	0.588	6082
ToxProblem=Freq	0.579	3961
ToxProblem=VeryFreq	0.588	410

Table 2: Statistically significant ($p < 0.05$) a posteriori unimodality results across all four datasets. Gray rows show negative values. Consult App. F for the full results.




Figure 6: Standard deviation of 30 randomly sampled (with replacement) pol_{obs} values (Eq. 5). The x-axis shows the sample size (number of annotators), starting from 3 and increasing one at a time, up to the largest annotator count for which a sufficient number of items exists (B).

Polarization attribution follows a monotonic trend Figure 5 shows *apunim* trends across ordinal sociodemographic attributes listed in Table 2. We find that polarization attribution changes consistently in one direction as we move from one end of an attribute scale to the other—for example, from irreligious to very religious annotators. This monotonic pattern appears across *all* ordinal sociodemographic dimensions and *all* datasets (apart from Age in the Kumar dataset, which is U-shaped), suggesting that perceived differences become stronger between groups that are further apart ideologically or demographically.

Finding 4. For ordinal sociodemographic dimensions, polarization attribution tends to grow stronger as the differences between annotator groups increase.

6 How many annotators do we need?

While *apunim* works best with many annotations per item, the cost required for multiple human annotators is often greatly restrictive (Rossi et al., 2024). Thus, it is imperative to find the least annotators needed for *apunim* to generate stable and reliable estimates. For each item per dataset, we randomly sample with replacement a subset of annotators and compute the standard deviation of 30 sampled pol_{obs} values (see Equation 5). Figure 6 demonstrates the effect of the number of annotators on the reliability of the polarization estimation, which we discuss next.

Each dataset features a different standard deviation even with the same number of annotators, which is broadly consistent with the amount of polarization in each dataset (App. B). However, we observe a clear, non-linear, inverse relationship between the number of annotators and the standard error across datasets. This can be attributed to a

more robust histogram estimation used by `nDFU` (§4.1) when provided with more points for each of the subgroups. In the case of our datasets, we observe that after about 20 points, the estimation is good enough that more annotations contribute trivially to the robustness of the metric.

7 Conclusion & Future work

We investigated whether fundamental differences can be attributed to annotator groups using polarization, instead of standard agreement measures. We discovered as of yet unknown issues regarding the presence of inherent polarization and conflicting polarization directions which we bypass using a new metric, *apunim*. We then applied our proposed metric to four hate-speech and toxicity detection datasets, discovering that (1) race and gender seem to generally drive polarization attribution, and (2) groups that are further apart increasingly drive polarization. Furthermore, we estimated the robustness of our metric w.r.t. the number of annotators, and released a PyPi installable library implementing it.

Our framework has broad applications across NLP tasks related to content moderation. Because many of these tasks involve multiple discussions, the framework can also be applied at the discussion level rather than only at the dataset level. This would allow researchers to identify particularly polarizing discussions and study the underlying sources of disagreement for specific minority groups. Discussions could also be analyzed by comment turn, enabling the study of how polarization develops over the course of an interaction. Finally, our approach can be used to more reliably evaluate whether LLM-as-a-judge systems reproduce the polarization patterns observed among different minority groups.

8 Limitations

Researchers often release only aggregated single-label annotations or omit annotator sociodemographic characteristics entirely. As a result, both our analysis and the applicability of our framework are limited by the scarcity of suitable datasets. In addition, because datasets with large numbers of annotators per item are rare, our approach operates only at the dataset level. This limitation prevents us from identifying individual comments or examples that could serve as concrete explanations for the observed *apunim* values. Finally, although we compare findings across datasets and relate them to prior work, we do not directly compare polarization and agreement metrics, since they rely on fundamentally different assumptions and mechanisms (see §2.2). This also makes it difficult to fully explain certain patterns in our results, such as why some groups show statistically significant

effects while others do not.

9 Ethical Considerations

All datasets used in this study were either publicly accessible, or were provided to us by the authors with full disclosure on the purpose of our research. No identifiable, personal information is included in our pipeline. We did not perform any annotation tasks ourselves.

References

- Sohail Akhtar, Valerio Basile, and Viviana Patti. 2019. A new measure of polarization in the annotation of hate speech. In *AI*IA 2019 – Advances in Artificial Intelligence*, pages 588–603, Cham. Springer International Publishing.
- Hala Al Kuwatly, Maximilian Wich, and Georg Groh. 2020. Identifying and measuring annotator bias based on annotators’ demographic characteristics. In *Proceedings of the fourth workshop on online abuse and harms*, pages 184–190.
- Lora Aroyo, Alex Taylor, Mark Díaz, Christopher Homan, Alicia Parrish, Gregory Serapio-García, Vinodkumar Prabhakaran, and Ding Wang. 2023. *Dices dataset: Diversity in conversational ai evaluation for safety*. In *Advances in Neural Information Processing Systems*, volume 36, pages 53330–53342. Curran Associates, Inc.
- Reuben Binns, Michael Veale, Max Van Kleek, and Nigel Shadbolt. 2017. Like trainer, like bot? inheritance of bias in algorithmic content moderation. In *International conference on social informatics*, pages 405–415. Springer.
- Ted Byrt, Janet Bishop, and John B. Carlin. 1993. *Bias, prevalence and kappa*. *Journal of Clinical Epidemiology*, 46(5):423–429.
- Alessandro Checco, A. Roitero, Eddy Maddalena, S. Mizzaro, and Gianluca Demartini. 2017. Let’s agree to disagree: Fixing agreement measures for crowdsourcing.
- S.Y. Chen, Z. Feng, and X. Yi. 2017. A general introduction to adjustment for multiple comparisons. *Journal of Thoracic Disease*, 9(6):1725–1729. Doi: 10.21037/jtd.2017.05.34; PMID: 28740688; PMCID: PMC5506159.
- Shelby Dioum. 2025. *What explains the increase in online hate speech?* Accessed: 2025-09-22.
- Alvan R. Feinstein and Domenic V. Cicchetti. 1990. *High agreement but low kappa: I. the problems of two paradoxes*. *Journal of Clinical Epidemiology*, 43(6):543–549.
- Sokratis-Vissarion Giannoutsos, Dimitris Tsirmpas, and John Pavlopoulos. 2025. Llm depolarization via aposteriori unimodality.

- Tommaso Giorgi, Lorenzo Cima, Tiziano Fagni, Marco Avvenuti, and Stefano Cresci. 2025. Human and llm biases in hate speech annotations: A socio-demographic analysis of annotators and targets. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 19, pages 653–670.
- Nitesh Goyal, Ian D Kivlichan, Rachel Rosen, and Lucy Vasserman. 2022. Is your toxicity my toxicity? exploring the impact of rater identity on toxicity annotation. *Proceedings of the ACM on Human-Computer Interaction*, 6(CSCW2):1–28.
- Sture Holm. 1979. [A simple sequentially rejective multiple test procedure](#). *Scandinavian Journal of Statistics*, 6(2):65–70.
- Jonathan Ivey, Susan Gauch, and David Jurgens. 2025. [NUTMEG: Separating signal from noise in annotator disagreement](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2874–2887, Suzhou, China. Association for Computational Linguistics.
- Petra Kralj Novak, Teresa Scantamburlo, Andraž Pelicon, Matteo Cinelli, Igor Mozetič, and Fabiana Zollo. 2022. Handling disagreement in hate speech modelling. In *Information Processing and Management of Uncertainty in Knowledge-Based Systems*, pages 681–695, Cham. Springer International Publishing.
- Deepak Kumar, Patrick Gage Kelley, Sunny Consolvo, Joshua Mason, Elie Bursztein, Zakir Durumeric, Kurt Thomas, and Michael Bailey. 2021. Designing toxic content classification for a diversity of perspectives. In *Proceedings of the Seventeenth USENIX Conference on Usable Privacy and Security*, SOUPS’21, USA. USENIX Association.
- Giuseppe Mignemi, Antonio Calcagnì, Andrea Spoto, and Ioanna Manolopoulou. 2024. [Mixture polarization in inter-rater agreement analysis: a bayesian nonparametric index](#). *Statistical Methods & Applications*, 33(1):325–355.
- John Pavlopoulos and Aristidis Likas. 2023. [Distance from unimodality for the assessment of opinion polarization](#). *Cognitive Computation*, 15(2):731–738.
- John Pavlopoulos and Aristidis Likas. 2024. [Polarized opinion detection improves the detection of toxic language](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1946–1958, St. Julian’s, Malta. Association for Computational Linguistics.
- Anne Quaranto. 2022. Dog whistles, covertly coded speech, and the practices that enable them. *Synthese*, 200(4):330.
- Luca Rossi, Katherine Harrison, and Irina Shklovski. 2024. [The problems of llm-generated data in social science research](#). *Sociologica*, 18(2):145–168.
- Pratik S Sachdeva, Renata Barreto, Claudia von Vacano, and Chris J Kennedy. 2022. Assessing annotator identity sensitivity via item response theory: A case study in a hate speech corpus. In *Proceedings of the 2022 ACM conference on fairness, accountability, and transparency*, pages 1585–1603.
- Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. 2019. [The risk of racial bias in hate speech detection](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1668–1678, Florence, Italy. Association for Computational Linguistics.
- Maarten Sap, Swabha Swayamdipta, Laura Vianna, Xuhui Zhou, Yejin Choi, and Noah A. Smith. 2022. [Annotators with attitudes: How annotator beliefs and identities bias toxic language detection](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5884–5906, Seattle, United States. Association for Computational Linguistics.
- Ole Tange. 2025. [Gnu parallel 20250822](#). GNU Parallel is a general parallelizer to run multiple serial command line programs in parallel without changing them.
- United Nations. 2025. [Targets of hate](#). Accessed: 2025-09-22.
- Mariken A. C. G. van der Velden, Felicia Loecherbach, Wouter van Atteveltdt, Antske Fokkens, Myrthe Reuver, and Kasper Welbers. 2025. [Whose truth is it anyway? an experiment on annotation bias in times of factual opinion polarization](#). *Communication Methods and Measures*, 19(4):332–349.
- Michael Wojatzki, Tobias Horsmann, Darina Gold, and Torsten Zesch. 2018. Do women perceive hate differently: Examining the relationship between hate speech, gender, and agreement judgments.
- Ka Wong, Praveen Paritosh, and Lora Aroyo. 2021. [Cross-replication reliability - an empirical approach to interpreting inter-rater reliability](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7053–7065, Online. Association for Computational Linguistics.

A Acronyms

LLM	Large Language Model
nDFU	normalized Distance From Unimodality
FWER	Family-Wise Error Rate
NLP	Natural Language Processing
GPLv3	GNU General Public Licence v3
AA	African American

B Annotation details

The number of annotators per comment for all datasets presented in §5.1 can be found in Figure 8, and descriptive statistics can be found in Table 3. The DICES-990 dataset has exactly 123 annotations per comment, while the other datasets exhibit a few variations around the central mean reported in Table 1. Interestingly, the sample used from the Kumar dataset exhibits a few comments that vastly exceed the promised 5 number of annotators. We included these comments for the experiment presented in Table 2, but not for the annotator number experiment in Figure 6 since they would produce uninterpretable figures.

C P-value estimation implementation

Since our test is parametric, it assumes that there exist (1) enough data-points to reliably run the means-test, (2) enough annotations per sociodemographic group to estimate nDFUs in Equations 3-5, (3) the apunims of the random partitions are independent and identically distributed. Condition (1) can be safely assumed for most datasets, because they feature enough items to reliably assume normal residuals (see §5.1). We discuss condition (2) in §6. Condition (3) is partly fulfilled; the partitions are identically distributed, but may be dependent on the sample distribution.

Replacing the parametric means test with a non-parametric alternative would relax assumptions about the underlying distribution but would not resolve the potential dependence between the random apunims. We initially attempted to implement a permutation test on the empirical (random apunim) distribution, but this resulted in excessive computational costs.

While the test operates on a single group in principle, in practice we often want to test polarization for all subgroups of a specific dimension. Since we are simultaneously considering multiple hypotheses, we apply a multiple comparison correction to the resulting p-values—specifically the Holms method (Holm, 1979). The test is parameterized by the Family-Wise Error Rate (FWER),

which is used to tune the strength of the correction; we can increase this value to make our test more conservative towards multiple hypotheses (Chen et al., 2017)). In general, it is safe to set FWER equal to the significance level of our test (e.g., $FWER = 0.95$ if we are looking for $p < 0.05$).

D Software

D.1 Installation and usage

We release the `apunim` package, which is a Python library that implements both the nDFU and apunim calculations, as well as its statistical test. The software is available on PyPi. We also provide high-level documentation in HTML format that includes an introduction, guide, and full low-level documentation for each of the two provided functions (<https://anonymous.4open.science/r/apunim-page/>).

The code is open-source, licensed under the GNU General Public Licence v3 (GPLv3) and can be found on the project’s repository (<https://anonymous.4open.science/r/aposteriori-unimodality-E8C1/>). The library is cross-platform, works for any Python 3 version, and requires very few dependencies (namely `numpy`, `scipy`, `statsmodels`).

D.2 Implementation details

By default, the computation of the apunim values happens on a single thread. In cases of serious computational cost, we recommend running the apunim calculation on multiple concurrent processes, since all operations are CPU-bound and thread-safe. This may be useful in cases where there are multiple sociodemographic dimensions, each of which is computationally expensive.

In order to lower the computational cost required, we re-use the random partitions ($\tilde{r}_i$), which are originally calculated to obtain the apriori polarization values (see Equation 3), for the p-value calculations (§4.3). This means that in our current implementation, the number of groups for estimating apriori polarization ($\tilde{r}_i$ in Equation 3), and the number of apunim values used in the p-value calculation have to be the same.

E Replication Details

We use version 1.0.1 of the `apunim` library for our experiments. We use 100 random iterations for the apriori value calculation (Equation 3) and the p-value estimation (§4.3), and set a FWER rate of 0.95.

The current version of the library handles large number of annotators very well due to the use of vectorized operations, but struggles when faced with a lot of comments when the annotators belong to multiple groups. Indeed, while the experiments for all datasets presented in §5.1 need no




Figure 7: Histogram of dataset polarization for the human datasets considered in this paper per sociodemographic dimension.

Table 3: Descriptive statistics for the number of annotations per comment, grouped by dataset.

	count	mean	std	min	25%	50%	75%	max
DICES-350	350	123	0	123	123	123	123	123
DICES-990	990	72.8313	1.1652	69	72	73	74	76
Kumar et al. 2021	20000	5.0925	4.8872	5	5	5	5	650
Sap et al. 2022	585	5.6359	0.7713	3	6	6	6	12

more than a minute to run, the Kumar dataset requires about 12 hours of computation even with a 20% sample. The annotator variance analysis presented in §6 also required more than 10 hours of computation, although that is expected due to the nature of the experiment, which is very computationally demanding, and a non-standard use of our library.

This bottleneck is likely caused by repeated creation of numpy arrays, resulting in repeated memory allocations for each subgroup, in each comment, for each sociodemographic dimension. We plan on profiling and addressing it in the next releases of our software. Until then, we suggest either using multiprocessing for each sociodemographic subgroup, or running the different experiments in

parallel (indeed, we use the `gnu-parallel` linux package (Tange, 2025) for our own experiments).

In Figure 6, we removed values where there was a lack of items annotated by the specified number of annotators (e.g., only 7% of comments in DICES-990 were annotated by more than 73 annotators—see B).

F Full Results




Figure 8: We calculate the ratio of comments for each dataset (y-axis) that had x number of annotators (x-axis). We truncate the x-axis to 130 annotators, disregarding very few comments in the Kumar dataset exceeding this threshold (see Table 3).

Group	apunim	support
Sap		
Age=GenX+	0.0192	743
Age=Millennial	0.0189	1792
Age=GenZ	-0.0153	239
Eth=Black	-0.2967	1041
Eth=White	0.1400	2048
Gender=Male	0.0942	1635
Gender=Female	-0.1810	1333

Table 4: Full apunim results for the Sap Dataset.

Group	apunim	support
Kumar		
Age=1) 18-24	-0.009	9742
Age=2) 25-34	0.032***	33203
Age=3) 35-44	-0.009	21055
Age=4) 45-54	0.001	11025
Age=5) 55-64	-0.004	6123
Age=6) 65+	0.000	2567
Edu=1) Doctoral degree	0.005	1019
Edu=2) Prof. degree	0.004	1321
Edu=3) Master's degree	0.015**	13420
Edu=4) Bachelor's degree	0.040***	34085
Edu=5) Associate degree	-0.011**	9122
Edu=6) College, no degree	-0.034***	16666
Edu=7) High School	-0.008*	7326
Edu=8) No high school	0.000	462
Eth=Asian	-0.002	5062
Eth=Black	0.011*	10999
Eth=Hispanic	-0.003	2389
Eth=Multiracial	-0.004	3682
Eth=Other	0.002	1507
Eth=Unknown	-0.003	1160
Eth=White	-0.038***	45306
Gender=Female	-0.028***	40686
Gender=Male	0.028***	37929
Gender=Nonbinary	-0.001	489
Gender=Unknown	-0.003	621
Target=False	-0.063***	45234
Target=True	0.035***	24491
Parent=No	-0.071***	37231
Parent=Prefer not to say	-0.002	831
Parent=Yes	0.054***	41493
Trans=No	-0.089***	11978
Trans=Prefer not to say	0.000	856
Trans=Yes	0.001	2306
Pol=Conservative	0.033***	23010
Pol=Independent	-0.016***	22818
Pol=Liberal	-0.029***	32917
Pol=Other	-0.005	1822
Pol=Prefer not to say	0.002	3548
Rel=1) No	-0.080***	26008
Rel=2) Not very	-0.015***	10055
Rel=3) Somewhat	0.021***	20169
Rel=4) Very	0.049***	27284
Rel=Prefer not to say	0.000	1179

Table 5: Full apunim results for the Kumar Dataset.




Figure 9: Apunim values across ordinal factor levels with at least two statistically significant values, expanded for each ordinal sociodemographic dimension. This figure is complementary to Figure 5. The values in the graphs were extracted from the Tables of F.

Group	apunim	support
Kumar		
Tox=False	0.031***	20185
Tox=True	-0.072***	43265
SO=Bisexual	0.018***	10005
SO=Heterosexual	-0.057***	38048
SO=Homosexual	-0.002	2975
SO=Other	-0.000	728
SO=Prefer not to say	-0.002	1751
Tech=1) Very Low	-0.001	784
Tech=2) Low	-0.008	7973
Tech=3) Mid	0.003	10720
Tech=4) High	-0.027***	40477
Tech=5) Very high	0.021***	22471
ToxP=1) Never	0.006	4962
ToxP=2) Rarely	0.018***	14457
ToxP=3) Occasionally	-0.006	29476
ToxP=4) Frequently	-0.022***	24774
ToxP=5) Very Frequently	-0.015***	10656

Table 6: Full apunim results for the Kumar Dataset (cont’ed).

Group	apunim	support
DICES-350		
Race=AfricanAm.	-0.0428	9077
Race=Asian	0.0659	8138
Age=Millenial	-0.0059	11268
Age=GenZ	-0.0066	17528
Edu=College or higher	0.0124	23475
Edu=High school or lower	-0.0152	12833
Edu=Other	-0.0214*	2191
Gender=Man	-0.0193	19093
Gender=Woman	0.0215	19406
Race=Latino	-0.0030	6886
Race=Multiracial	-0.0001	5008
Race=White	-0.0198	9390

Table 7: Full apunim results for the DICES-350 Dataset.

Group	apunim	support
DICES-990		
Age=GenX+	-0.0471	14384
Age=Millenial	-0.0043	38436
Age=GenZ	0.0644	13741
Edu=College or higher	0.0095	59142
Edu=High school or lower	-0.0751	7419
Gender=Man	0.0278	32494
Gender=Woman	-0.0235	33471
Race=AfricanAm.	0.0909	5810
Race=Latino	-0.0030	6886
Race=Multiracial	-0.0001	5008
Race=White	-0.0198	9390

Table 8: Full apunim results for the DICES-990 Dataset.