

Explanation of Revisions

The previous version of this manuscript was desk rejected due to formatting issues before reviewer feedback was provided. We nevertheless outline the changes made to the manuscript since. All references to Figures, Tables and Sections refer to the numbers reported in the current submission, unless otherwise stated.

Formatting

A formatting oversight on our part led to the accidental adoption of 10pt fontsize instead of the ACL standard of 11pt. The error has been corrected.

Major changes

1. We added standard deviation to the experiment presented in (current) Figure 6 instead of just point estimates.
2. We added a brief discussion section around current NLP datasets and their interaction with our proposed methodology to better frame our call for smaller datasets with more annotators per item.
3. The Equation of the apunim metric (Eq.8) was accidentally implemented with the wrong sign. The sign was inverted and the experiments re-ran. We also fixed minor issues with dataset filtering, which slightly impacted results.
4. We slightly refactored Finding 4. Specifically, we now focus more on the impact of Religion specifically as well as the dichotomy between ideological and innate personal characteristics.
5. We added computational and methodological limitations in the Limitations and dual-use concerns in the Ethical Considerations.

Supplementary Experiments

1. Information about the synthetic experiments can now be found in Appendix F.

2. A brief qualitative analysis on comments exhibiting Inherent Polarization was conducted in Appendix G.
3. We added ablations for both sample-sizes and sample selection for the Kumar dataset (Appendix H).

Improvements

1. We spit off some subsections into their own sections in order to improve readability (specifically, current Sections 3 and 7, Appendix E).
2. Clarified the hypotheses presented in Section 4.2.4 and explicitly singled-out the DICES datasets.
3. Fixed minor issues with Table 1, and improved visual clarity in Table 2.
4. Figure 5 now also includes statistical significance. Added a footnote explaining statistical insignificance around zero-values.
5. Appendix B now reports granular information about annotation counts per dataset and examined groups within each Personal Characteristic dimension.
6. Split Limitations and Ethical Considerations in paragraphs to increase readability.